I revised venus-ravi2:/usr/local/ravi/post_proc/check_header.c to print out the first few values from a data file along with a header. From

./check_header -l localhost /mnt/PESD_data/vsr1a.1w1.10-113-2200.raw

I got this for the data, in (real, imaginary) pairs:

```
 -1.5 -11.5
 -2.5  -7.5
-35.5  -8.5
 70.5   6.5
 33.5 -29.5
-54.5 -22.5
....
```

Reading the first few data bytes from that file into a buffer through an sshfs mount and then unpacking that with struct.unpack("!16b",buf[0:16]) and adding 0.5 I got:

```
-16.5, -11.5, -6.5,  -1.5,  16.5,  -7.5, -14.5,  -2.5,
-49.5,  -8.5,  6.5, -35.5,  34.5,   6.5,  41.5,  70.5
```

I've used boldface to show how they match up. Using struct.unpack("!xbxbxbxbxbxbxbxb",buf[0:16]) I get this:

```
 (-12, -2, -8, -3, -9, -36, 6, 70)
```

which, adding 0.5, gives what's expected. So where do these unexpected bytes come from? Doing this on ravi2:

```
kuiper@venus-ravi2:/usr/local/ravi/post_proc$ od -j 152 -N 16 -t d1
/mnt/PESD_data/vsr1a.1w1.10-113-2200.raw
0000230  -17  -12   -7   -2   16   -8  -15   -3  -50   -9    6  -36   34    6   41   70
```

which confirms what I see over an sshfs mount. Is it really possible that vdr_io between a 32-bit machine (venus-vsr1) and a 64-bit machine (ravi2) uses only every other byte on disk?